

Anomaly Detection in Transactions using Python

Anomaly detection in transactions means identifying unusual or unexpected patterns within transactions or related activities. These patterns, known as anomalies or outliers, deviate significantly from the expected norm and could indicate irregular or fraudulent behaviour. If you want to learn how to detect anomalies in transactions, this is for you.

Anomaly Detection in Transactions: Process We Can Follow

Anomaly detection plays a crucial role in various businesses, especially those dealing with financial transactions, online activities, and security-sensitive operations.

We can follow a systematic process to address the challenge of anomaly detection. We can begin by collecting and preparing transaction data, ensuring its accuracy and consistency. Then, we can find patterns in the data to find anomalies and use specialized anomaly detection algorithms like isolation forest to detect anomalies.

So the process starts with data collection. I have found an ideal dataset that can be used for detecting anomalies in transactions.

```
In [1]: import pandas as pd
import plotly.express as px
from sklearn.model_selection import train_test_split
from sklearn.ensemble import IsolationForest
from sklearn.metrics import classification_report

data = pd.read_csv("C:\\Users\\kourg\\Downloads\\transaction_anomalies_dataset.csv")
print(data.head())
```

	Transaction_ID	Transaction_Amount	Transaction_Volume	\
0	TX0	1024.835708	3	
1	TX1	1013.952065	4	
2	TX2	970.956093	1	
3	TX3	1040.822254	2	
4	TX4	998.777241	1	

	Average_Transaction_Amount	Frequency_of_Transactions	\
0	997.234714	12	
1	1020.210306	7	
2	989.496604	5	
3	969.522480	16	
4	1007.111026	7	

	Time_Since_Last_Transaction	Day_of_Week	Time_of_Day	Age	Gender	Income	\
0	29	Friday	06:00	36	Male	1436074	
1	22	Friday	01:00	41	Female	627069	
2	12	Tuesday	21:00	61	Male	786232	
3	28	Sunday	14:00	61	Male	619030	
4	7	Friday	08:00	56	Female	649457	

	Account_Type
0	Savings
1	Savings
2	Savings
3	Savings
4	Savings

```
In [2]: ▶ print(data.isnull().sum())
```

```
Transaction_ID          0
Transaction_Amount      0
Transaction_Volume      0
Average_Transaction_Amount  0
Frequency_of_Transactions  0
Time_Since_Last_Transaction  0
Day_of_Week             0
Time_of_Day             0
Age                     0
Gender                  0
Income                  0
Account_Type            0
dtype: int64
```

```
In [3]: ▶ print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Transaction_ID                        1000 non-null   object
 1   Transaction_Amount                    1000 non-null   float64
 2   Transaction_Volume                    1000 non-null   int64
 3   Average_Transaction_Amount            1000 non-null   float64
 4   Frequency_of_Transactions              1000 non-null   int64
 5   Time_Since_Last_Transaction            1000 non-null   int64
 6   Day_of_Week                           1000 non-null   object
 7   Time_of_Day                           1000 non-null   object
 8   Age                                    1000 non-null   int64
 9   Gender                                1000 non-null   object
10   Income                                1000 non-null   int64
11   Account_Type                           1000 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 93.9+ KB
None
```

In [4]: `print(data.describe())`


	Transaction_Amount	Transaction_Volume	Average_Transaction_Amount	\
count	1000.000000	1000.000000	1000.000000	
mean	1038.122511	2.498000	1000.682506	
std	283.580055	1.115006	20.632334	
min	849.024392	1.000000	939.081423	
25%	966.028796	1.000000	986.800556	
50%	1002.118678	3.000000	1000.501902	
75%	1033.143657	3.000000	1015.155595	
max	3227.459108	4.000000	1073.154036	

	Frequency_of_Transactions	Time_Since_Last_Transaction	Age	\
count	1000.000000	1000.000000	1000.000000	
mean	12.078000	15.341000	40.641000	
std	4.245225	8.361258	13.819953	
min	5.000000	1.000000	18.000000	
25%	8.000000	8.000000	29.000000	
50%	12.000000	16.000000	41.000000	
75%	16.000000	22.000000	53.000000	
max	19.000000	29.000000	64.000000	

	Income
count	1.000000e+03
mean	8.948238e+05
std	3.453562e+05
min	3.001590e+05
25%	5.917308e+05
50%	8.876645e+05
75%	1.178102e+06
max	1.499670e+06


```
In [5]: ▶ # Distribution of Transaction Amount
fig_amount = px.histogram(data, x='Transaction_Amount',
                           nbins=20,
                           title='Distribution of Transaction Amount')
fig_amount.show()
```

Distribution of Transaction Amount



```
# Transaction Amount by Account Type
fig_box_amount = px.box(data,
                        x='Account_Type',
                        y='Transaction_Amount',
                        title='Transaction Amount by Account Type')
fig_box_amount.show()
```


Transaction Amount by Account Type



```
In [7]: ▶ # Average Transaction Amount vs. Age
fig_scatter_avg_amount_age = px.scatter(data, x='Age',
                                         y='Average_Transaction_Amount',
                                         color='Account_Type',
                                         title='Average Transaction Amount vs. Age',
                                         trendline='ols')


fig_scatter_avg_amount_age.show()
```

Average Transaction Amount vs. Age



```
In [8]: ▶ # Count of Transactions by Day of the Week  
fig_day_of_week = px.bar(data, x='Day_of_Week',  
                          title='Count of Transactions by Day of the Week')  
fig_day_of_week.show()
```

Count of Transactions by Day of the Week

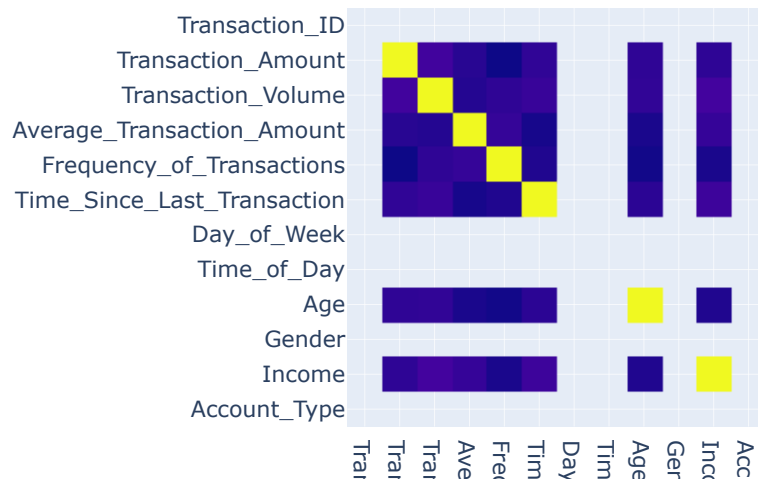


```
In [10]: ► import pandas as pd

# Assuming 'data' is your DataFrame
data_numeric = data.apply(pd.to_numeric, errors='coerce')

# Correlation Heatmap
correlation_matrix = data_numeric.corr()
fig_corr_heatmap = px.imshow(correlation_matrix, title='Correlation Heatmap')
fig_corr_heatmap.show()
```

Correlation Heatmap





```
In [11]: ▶ # Calculate mean and standard deviation of Transaction Amount
mean_amount = data['Transaction_Amount'].mean()
std_amount = data['Transaction_Amount'].std()

# Define the anomaly threshold (2 standard deviations from the mean)
anomaly_threshold = mean_amount + 2 * std_amount

# Flag anomalies
data['Is_Anomaly'] = data['Transaction_Amount'] > anomaly_threshold

# Scatter plot of Transaction Amount with anomalies highlighted
fig_anomalies = px.scatter(data, x='Transaction_Amount', y='Average_Transaction_Amount',
                           color='Is_Anomaly', title='Anomalies in Transaction Amount')
fig_anomalies.update_traces(marker=dict(size=12),
                             selector=dict(mode='markers', marker_size=1))
fig_anomalies.show()
```

Anomalies in Transaction Amount



```
In [12]: ► # Calculate the number of anomalies
num_anomalies = data['Is_Anomaly'].sum()

# Calculate the total number of instances in the dataset
total_instances = data.shape[0]

# Calculate the ratio of anomalies
anomaly_ratio = num_anomalies / total_instances
print(anomaly_ratio)
```

0.02

```
In [14]: ► relevant_features = ['Transaction_Amount',
                                'Average_Transaction_Amount',
                                'Frequency_of_Transactions']

# Split data into features (X) and target variable (y)
X = data[relevant_features]
y = data['Is_Anomaly']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train the Isolation Forest model
model = IsolationForest(contamination=0.02, random_state=42)
model.fit(X_train)
```

Out[14]: ► IsolationForest

Here we are training an anomaly detection model using the Isolation Forest algorithm. First, we selected the relevant features for detection, namely Transaction_Amount, Average_Transaction_Amount, and Frequency_of_Transactions.

We split the dataset into features (X) and the target variable (y), where X contains the selected features and y contains the binary labels indicating whether an instance is an anomaly or not. Then, we further split the data into training and testing sets using an 80-20 split ratio. Next, we created an Isolation Forest model with a specified contamination parameter of 0.02 (indicating the expected ratio of anomalies) and a random seed for reproducibility. The model is then trained on the training set (X_train).

Now let's have a look at the performance of this anomaly detection model:

```
In [15]: ▶ # Predict anomalies on the test set
y_pred = model.predict(X_test)

# Convert predictions to binary values (0: normal, 1: anomaly)
y_pred_binary = [1 if pred == -1 else 0 for pred in y_pred]

# Evaluate the model's performance
report = classification_report(y_test, y_pred_binary, target_names=['Normal', 'Anomaly'])
print(report)
```

	precision	recall	f1-score	support
Normal	1.00	1.00	1.00	196
Anomaly	1.00	1.00	1.00	4
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

```
In [16]: ▶ # Relevant features used during training
relevant_features = ['Transaction_Amount', 'Average_Transaction_Amount', 'Frequency_of_Transactions']

# Get user inputs for features
user_inputs = []
for feature in relevant_features:
    user_input = float(input(f"Enter the value for '{feature}': "))
    user_inputs.append(user_input)

# Create a DataFrame from user inputs
user_df = pd.DataFrame([user_inputs], columns=relevant_features)

# Predict anomalies using the model
user_anomaly_pred = model.predict(user_df)

# Convert the prediction to binary value (0: normal, 1: anomaly)
user_anomaly_pred_binary = 1 if user_anomaly_pred == -1 else 0

if user_anomaly_pred_binary == 1:
    print("Anomaly detected: This transaction is flagged as an anomaly.")
else:
    print("No anomaly detected: This transaction is normal.")
```

```
Enter the value for 'Transaction_Amount': 5889
Enter the value for 'Average_Transaction_Amount': 3000
Enter the value for 'Frequency_of_Transactions': 5
Anomaly detected: This transaction is flagged as an anomaly.
```

SUMMARY

So this is how you can perform anomaly detection in transactions using Machine Learning and Python. Anomaly detection in transactions means identifying unusual or unexpected patterns within transactions or related activities. These patterns, known as anomalies or outliers, deviate significantly from the expected norm and could indicate irregular or fraudulent behaviour. I hope you liked this article on Anomaly Detection in transactions using Python.

